

# MSDS Practicum I - HR analysis

March 1, 2020

## 0.1 Sylvester Brown

## 0.2 Regis University

## 0.3 Praticum I

### 0.3.1 Spring 2020

### 0.3.2 Purpose of project:

Showcase skills in data manipulation and engineering, exploratory data analysis, visualizations and Machine leaning.

### 0.3.3 The problems to solve are:

1. what main characteristics contribute to the reason of why employees are leaving?
2. Which learning model appears to be better for predicting which employees will leave.

### 0.3.4 About the Dataset:

The dataset belongs to William Walter. Data can be found at “<https://www.kaggle.com/colara/human-resource>”

### 0.3.5 Other Resources used:

”Pandas for Everyone’ by Daniel Chen

<https://www.kaggle.com/colara/human-resources-analytics-a-descriptive-analysis>

<https://www.kaggle.com/daphnecor/predict-employee-turnover-rate-0>

<https://www.kaggle.com/henryshtang/hr-data-exploration>

<https://www.kaggle.com/rhuebner/human-resources-data-set/kernels>

### 0.3.6 Import Libraries

```
[1]: # import numpy, pandas, seaborn and matplotlib

import numpy as np
import pandas as pd
import sklearn
import seaborn as sns
```

```
sns.set()
import matplotlib.pyplot as plt
%matplotlib inline
# import warnings filter
from warnings import simplefilter
# ignore all future warnings
simplefilter(action='ignore', category=FutureWarning)
```

### 0.3.7 Upload Data

```
[2]: # Use pandas to read the csv file

hr = pd.read_csv('C:/Users/spbro/OneDrive/Desktop/Human Resources.csv')
```

### 0.3.8 Exploratory Data Analysis

```
[3]: # List the column names for the dataset
hr.columns
```

```
[3]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
          'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
          'promotion_last_5years', 'sales', 'salary'],
          dtype='object')
```

A closer look at the column names shows that some of the columns are not descriptive enough to help the analyst know what the column contains. For this reason the “Sales” column will need to be changed to “departments” and “average\_monthly\_hours” will be changed to “average\_monthly\_hours”. “Work\_accidents” change to “work\_accidents”, “time\_spend\_company” to “time\_spent\_at\_company”, “number\_project” to “number\_of\_projects”

```
[4]: # use rename() to change column names

hr = hr.rename(columns = {'sales':'department'})
hr = hr.rename(columns = {'average_monthly_hours':'average_monthly_hours'})
hr = hr.rename(columns = {'Work_accidents':'work_accident'})
hr = hr.rename(columns = {'time_spend_company':'time_spent_at_company'})
hr = hr.rename(columns = {'number_project':'number_of_projects'})
```

```
[5]: # Display first 5 rows
```

```
hr.head()
```

```
[5]:
```

	satisfaction_level	last_evaluation	number_of_projects	\
0	0.38	0.53	2	
1	0.80	0.86	5	
2	0.11	0.88	7	
3	0.72	0.87	5	

4	0.37	0.52	2	
---	------	------	---	--

	average_monthly_hours	time_spent_at_company	Work_accident	left	\
0	157	3	0	1	
1	262	6	0	1	
2	272	4	0	1	
3	223	5	0	1	
4	159	3	0	1	

	promotion_last_5years	department	salary
0	0	sales	low
1	0	sales	medium
2	0	sales	medium
3	0	sales	low
4	0	sales	low

```
[6]: # Display last 5 rows
```

```
hr.tail()
```

```
[6]:
```

	satisfaction_level	last_evaluation	number_of_projects	\
14994	0.40	0.57	2	
14995	0.37	0.48	2	
14996	0.37	0.53	2	
14997	0.11	0.96	6	
14998	0.37	0.52	2	

	average_monthly_hours	time_spent_at_company	Work_accident	left	\
14994	151	3	0	1	
14995	160	3	0	1	
14996	143	3	0	1	
14997	280	4	0	1	
14998	158	3	0	1	

	promotion_last_5years	department	salary
14994	0	support	low
14995	0	support	low
14996	0	support	low
14997	0	support	low
14998	0	support	low

```
[7]: hr.shape
```

```
# There are 14999 rows and 10 columns
```

```
[7]: (14999, 10)
```

```
[8]: # A statistical description of each column
```

```
hr.describe()
```

```
[8]:
```

	satisfaction_level	last_evaluation	number_of_projects	\
count	14999.000000	14999.000000	14999.000000	
mean	0.612834	0.716102	3.803054	
std	0.248631	0.171169	1.232592	
min	0.090000	0.360000	2.000000	
25%	0.440000	0.560000	3.000000	
50%	0.640000	0.720000	4.000000	
75%	0.820000	0.870000	5.000000	
max	1.000000	1.000000	7.000000	

	average_monthly_hours	time_spent_at_company	Work_accident	\
count	14999.000000	14999.000000	14999.000000	
mean	201.050337	3.498233	0.144610	
std	49.943099	1.460136	0.351719	
min	96.000000	2.000000	0.000000	
25%	156.000000	3.000000	0.000000	
50%	200.000000	3.000000	0.000000	
75%	245.000000	4.000000	0.000000	
max	310.000000	10.000000	1.000000	

	left	promotion_last_5years
count	14999.000000	14999.000000
mean	0.238083	0.021268
std	0.425924	0.144281
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

```
[9]: # check if data contains any null values
```

```
hr.isnull().any()
```

```
[9]:
```

satisfaction_level	False
last_evaluation	False
number_of_projects	False
average_monthly_hours	False
time_spent_at_company	False
Work_accident	False
left	False
promotion_last_5years	False
department	False

```
salary                False
dtype: bool
```

```
[10]: # display more descriptive stats using .info()
hr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
satisfaction_level    14999 non-null float64
last_evaluation       14999 non-null float64
number_of_projects    14999 non-null int64
average_monthly_hours 14999 non-null int64
time_spent_at_company  14999 non-null int64
Work_accident         14999 non-null int64
left                 14999 non-null int64
promotion_last_5years 14999 non-null int64
department            14999 non-null object
salary               14999 non-null object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

```
[11]: # display unique values of hr['salary']

hr['salary'].unique()
```

```
[11]: array(['low', 'medium', 'high'], dtype=object)
```

```
[12]: 5# Convert salary variable type to numeric
hr['salary'] = hr['salary'].map({'low':1, 'medium':2, 'high':3})

hr['salary'].unique()
```

```
[12]: array([1, 2, 3], dtype=int64)
```

```
[13]: # display unique values of hr['department']
hr['department'].unique()
```

```
[13]: array(['sales', 'accounting', 'hr', 'technical', 'support', 'management',
            'IT', 'product_mng', 'marketing', 'RandD'], dtype=object)
```

```
[14]: #numpy.where(condition[, x, y])
      #Return elements, either from x or y, depending on condition.

      # change 'support' category to 'technical' category
hr['department'] = np.where(hr['department'] == 'support', 'technical',
                             ↪hr['department'])
```

```
# change 'IT' in 'technical' category'
hr['department'] = np.where(hr['department'] == 'IT' , 'technical',
    ↪hr['department'])

hr['department'].unique()
```

```
[14]: array(['sales', 'accounting', 'hr', 'technical', 'management',
        'product_mng', 'marketing', 'RandD'], dtype=object)
```

explore the number of employees that left the company versus those that did not leave using count and groupby

```
[15]: # how many people left the company using count

hr['left'].value_counts()
```

```
[15]: 0    11428
      1     3571
      Name: left, dtype: int64
```

```
[16]: # (3571/14999)*100 = 23.808% This indicated that about 24% of the employees
      # left the company
```

```
[17]: # Use groupby to explore those who left per each cloumn using mean values

hr.groupby('left').mean()
```

```
[17]:      satisfaction_level  last_evaluation  number_of_projects \
left
0                0.666810           0.715473           3.786664
1                0.440098           0.718113           3.855503

      average_monthly_hours  time_spent_at_company  Work_accident \
left
0                199.060203           3.380032           0.175009
1                207.419210           3.876505           0.047326

      promotion_last_5years      salary
left
0                0.026251  1.650945
1                0.005321  1.414730
```

## 0.4 Observations

- The average satisfaction level of employees who stayed with the company is higher (66%) than that of the employees who left(44%).

- The average monthly work hours of employees who left the company is more than that of the employees who stayed. The average hours worked per month of those who left the company was 207.4hrs and those who stayed was about 199hrs.
- The employees who had workplace accidents are less likely to leave than that of the employee who did not have workplace accidents.
- The employees who were promoted in the last five years are less likely to leave than those who did not get a promotion in the last five years.
- On face value, it appears there was no significant different in the mean between those who left and those who did not leave when it comes to 'last\_evaluations' and 'number of projects'
- The average number of people who stayed because of salary was higher than those whose left.

```
[18]: # Use groupby to explore the data using department
```

```
hr.groupby('department').mean()
```

```
[18]:
```

	satisfaction_level	last_evaluation	number_of_projects	\
department				
RandD	0.619822	0.712122	3.853875	
accounting	0.582151	0.717718	3.825293	
hr	0.598809	0.708850	3.654939	
management	0.621349	0.724000	3.860317	
marketing	0.618601	0.715886	3.687646	
product_mng	0.619634	0.714756	3.807095	
sales	0.614447	0.709717	3.776329	
technical	0.613687	0.720976	3.839054	

	average_monthly_hours	time_spent_at_company	Work_accident	\
department				
RandD	200.800508	3.367217	0.170267	
accounting	201.162973	3.522816	0.125163	
hr	198.684709	3.355886	0.120433	
management	201.249206	4.303175	0.163492	
marketing	199.385781	3.569930	0.160839	
product_mng	199.965632	3.475610	0.146341	
sales	200.911353	3.534058	0.141787	
technical	201.813795	3.416127	0.144106	

	left	promotion_last_5years	salary
department			
RandD	0.153748	0.034307	1.602287
accounting	0.265971	0.018253	1.629726
hr	0.290934	0.020298	1.607578
management	0.144444	0.109524	2.071429
marketing	0.236597	0.050117	1.624709
product_mng	0.219512	0.000000	1.575388

sales	0.244928	0.024155	1.557971
technical	0.246924	0.008258	1.562500

## 0.5 Observations

- Management department had the most satisfaction level
- The department with the highest average of people who left was hr and management had the least.
- the satisfaction level for hr was about 59%. Accounting has the lowest average satisfaction level at about 58%.
- It can be observed that satisfaction level may not be the only criteria for leaving, given that accounting did not become the department with the most number of people who left seeing that their satisfaction level was the lowest in the company. In other words, other tools like correlation analysis may have to be used to estimate what combinations of factors may be a better predictor of does who left the company.

```
[19]: # Use groupby to explore the data using salary
```

```
hr.groupby('salary').mean()
```

```
[19]:
```

	satisfaction_level	last_evaluation	number_of_projects	\
salary				
1	0.600753	0.717017	3.799891	
2	0.621817	0.717322	3.813528	
3	0.637470	0.704325	3.767179	

	average_monthly_hours	time_spent_at_company	Work_accident	left	\
salary					
1	200.996583	3.438218	0.142154	0.296884	
2	201.338349	3.529010	0.145361	0.204313	
3	199.867421	3.692805	0.155214	0.066289	

	promotion_last_5years
salary	
1	0.009021
2	0.028079
3	0.058205

## 0.6 Observations

- Given that 1 is low salary and 3 is for high salary, it appears that those with the high salary are usually the most satisfied as well.
- Those with the high salary appear to also have the lowest average of people who left.



## 0.7 Additional Questions that can be answered through more EDA

[20]: *# How many employees work in each department?*

```
hr['department'].value_counts()
```

```
[20]: technical      6176
      sales         4140
      product_mng   902
      marketing     858
      RandD         787
      accounting    767
      hr            739
      management    630
      Name: department, dtype: int64
```

[21]: *# How many employees per Salary range?*

```
hr['salary'].value_counts()
```

```
[21]: 1    7316
      2    6446
      3    1237
      Name: salary, dtype: int64
```

[22]: *# How many employees per salary range and department?*

```
table = hr.pivot_table(values="satisfaction_level", index="department",  
→columns="salary",aggfunc=np.count_nonzero)
```

```
table
```

```
[22]: salary      1      2      3
      department
      RandD      364.0  372.0  51.0
      accounting  358.0  335.0  74.0
      hr          335.0  359.0  45.0
      management  180.0  225.0  225.0
      marketing   402.0  376.0  80.0
      product_mng 451.0  383.0  68.0
      sales       2099.0 1772.0 269.0
      technical   3127.0 2624.0 425.0
```

## 0.8 Data Visualization

### 0.8.1 BOXPLOTS

```
[23]: f, axes = plt.subplots(3,2, figsize=(10,10), sharex=True)

plt.subplots_adjust(wspace=0.5) # adjust the space between the plots

sns.despine(left=True)

# plot a boxplot of satisfaction_level to see if there is outliers
sns.boxplot( x= 'satisfaction_level', data=hr, orient='v',ax=axes[0,0])

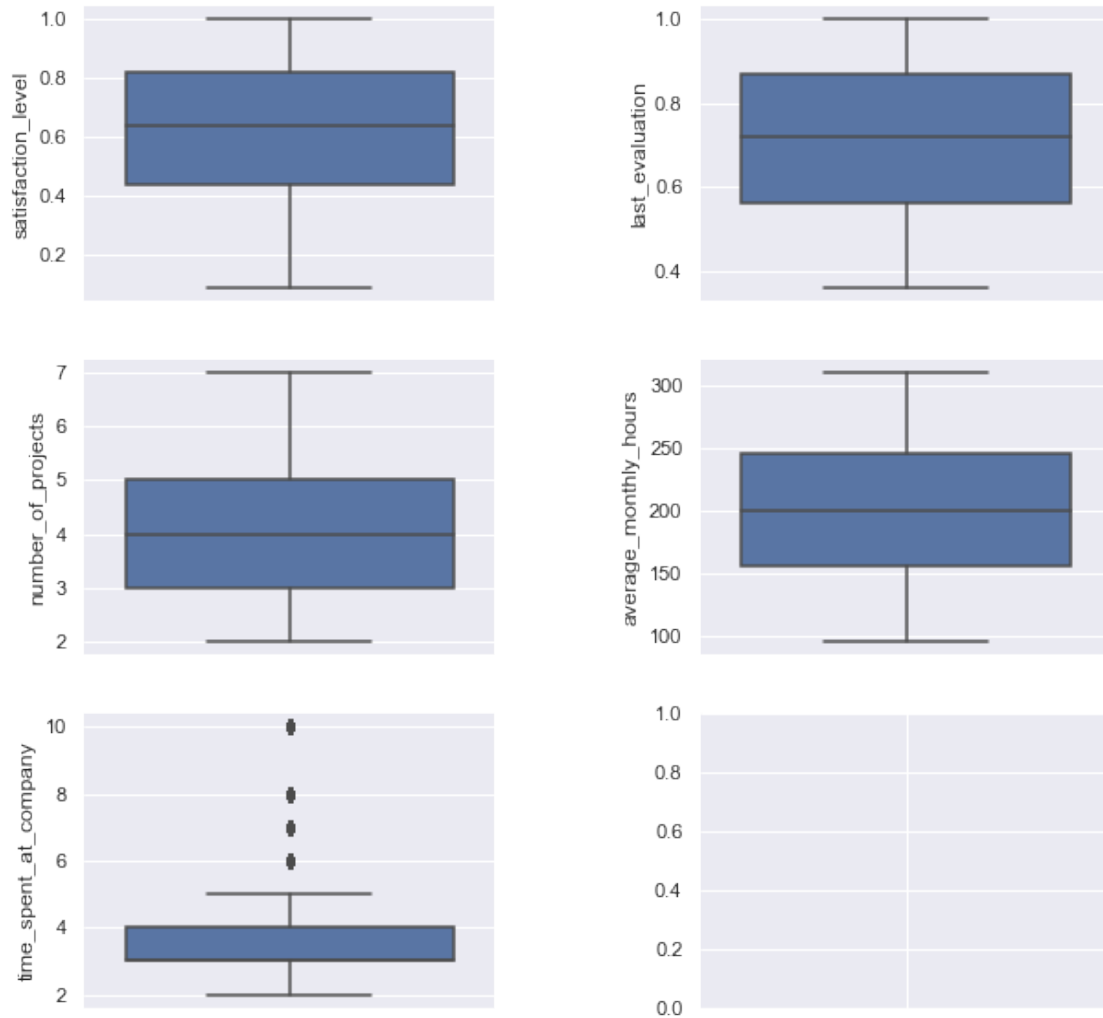
# plot a boxplot of last_evaluation to see if there is outliers
sns.boxplot( x= 'last_evaluation', data=hr, orient='v',ax=axes[0,1])

# plot a boxplot of number_project to see if there is outliers
sns.boxplot( x= 'number_of_projects', data=hr, orient='v',ax=axes[1,0])

# plot a boxplot of average_monthly_hours to see if there is outliers
sns.boxplot( x= 'average_monthly_hours', data=hr, orient='v',ax=axes[1,1])

# plot a boxplot of time_spent_at_company to see if there is outliers
sns.boxplot( x= 'time_spent_at_company', data=hr, orient='v',ax=axes[2,0])
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x1772a891748>
```



## 0.9 Observation

Time spent at the company has outliers. The company also appears to be a young company.

## 0.10 BAR PLOT

[24]: *# The plot shows the amount of employees that stayed and left the company.*


```
plt.figure(figsize=(4,5))
ax = sns.countplot(hr.left)
total = float(len(hr))

for p in ax.patches:
    height = p.get_height()
    ax.text(p.get_x()+p.get_width()/2.,
```

```

        height + 3,
        '{:1.2f}'.format(height/total),
        ha="center")
plt.title('Stayed or Left', fontsize=14);

```



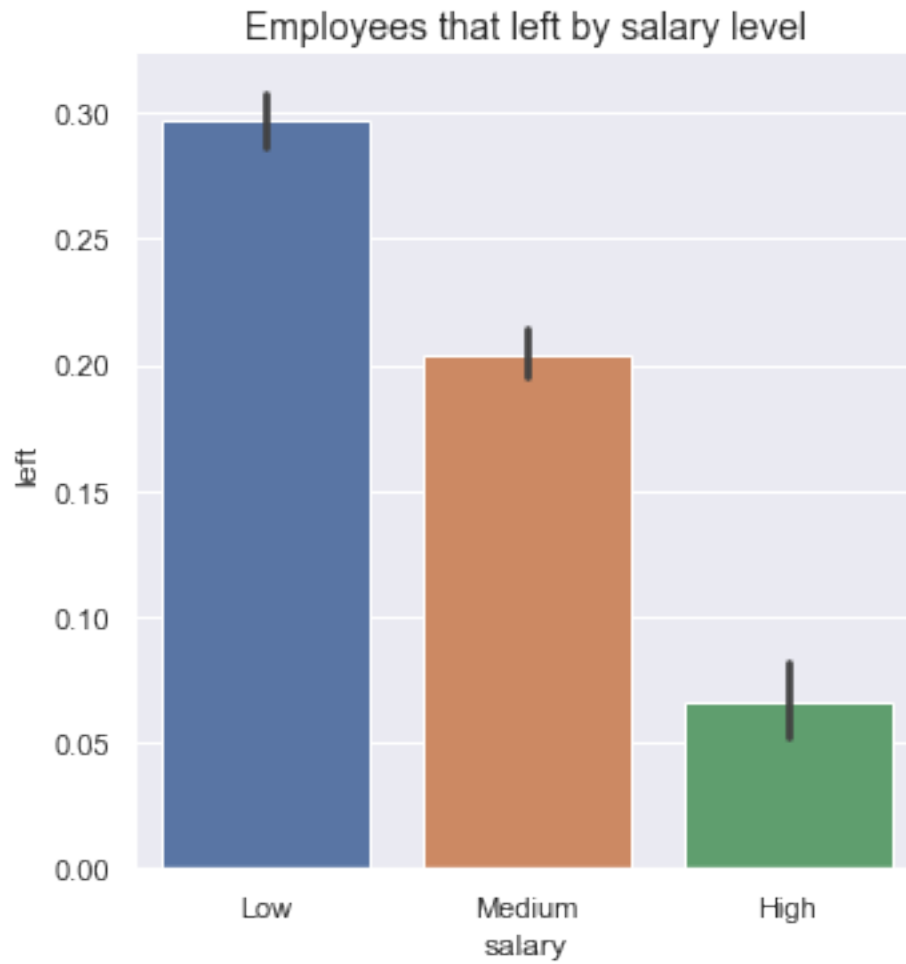
```
[25]: # plot of Employees that left by Salary level
```

```

j = sns.catplot(x='salary', y='left', kind='bar', data=hr)
plt.title('Employees that left by salary level', fontsize=14)
j.set_xticklabels(['Low', 'Medium', 'High'])


```

```
[25]: <seaborn.axisgrid.FacetGrid at 0x17728efebc8>
```



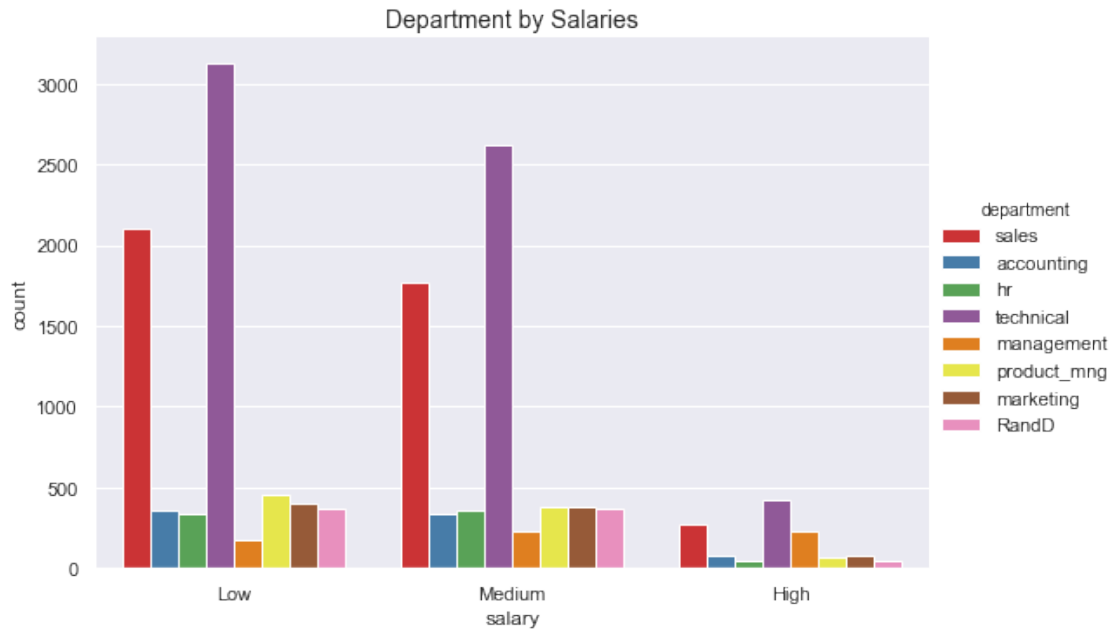
```
[26]: # plot of Employees that left by Salary level  
  
j = sns.catplot(x='Work_accident', y='left', kind='bar', data=hr)  
plt.title('Employees that left by work_accidents', fontsize=14)
```

```
[26]: Text(0.5, 1, 'Employees that left by work_accidents')
```



```
[27]: # plot of departments by salary level
```

```
h = sns.catplot(x = 'salary', hue='department', kind = 'count', height = 5, aspect=1.5, data=hr, palette='Set1' )  
plt.title("Department by Salaries", fontsize=14)  
h.set_xticklabels(['Low', 'Medium', 'High']);
```

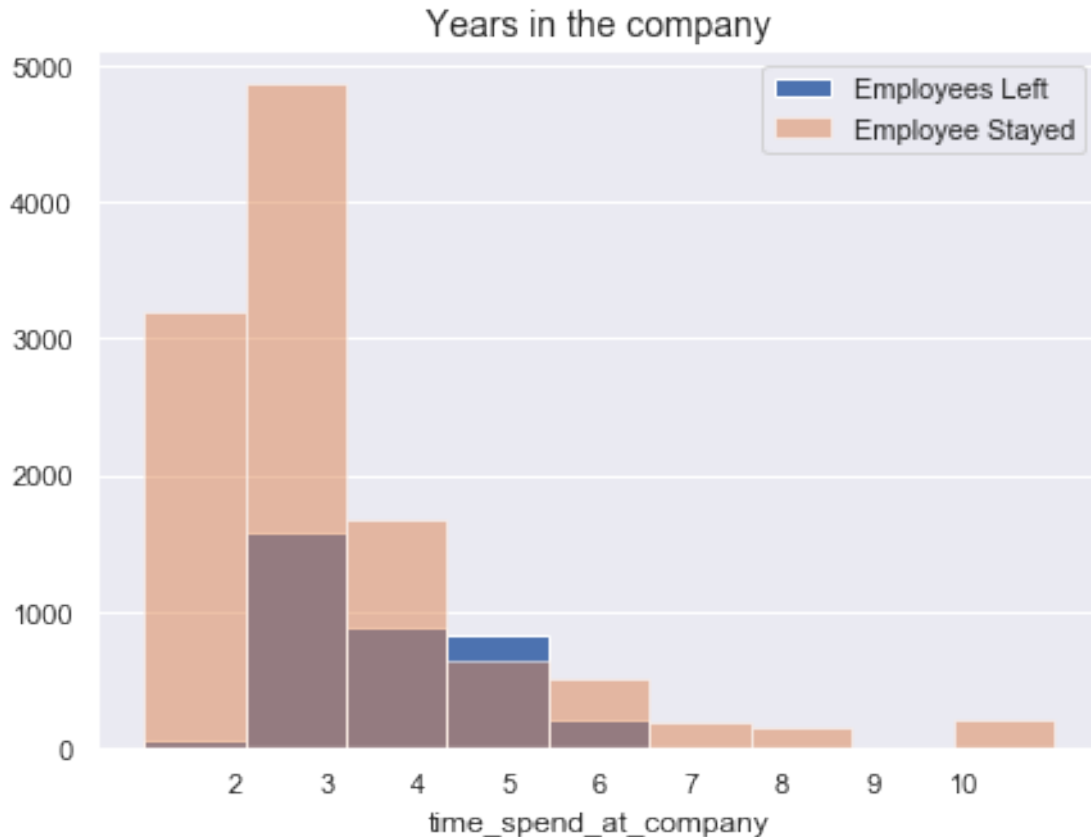


## 0.11 OBSERVATION

- Most of the employees of the technical department have low or medium salaries.
- Sales department is in the second place where most of the employees receives low and medium salaries.

```
[28]: # Plot comparing employees that left with years in company

plt.figure(figsize =(7,5))
bins = np.linspace(1.0, 11,10)
plt.hist(hr[hr['left']==1]['time_spent_at_company'], bins=bins, alpha=1,
        ↪label='Employees Left')
plt.hist(hr[hr['left']==0]['time_spent_at_company'], bins=bins, alpha = 0.5,
        ↪label = 'Employee Stayed')
plt.grid(axis='x')
plt.xticks(np.arange(2,11))
plt.xlabel('time_spend_at_company')
plt.title('Years in the company', fontsize=14)
plt.legend(loc='best');
```



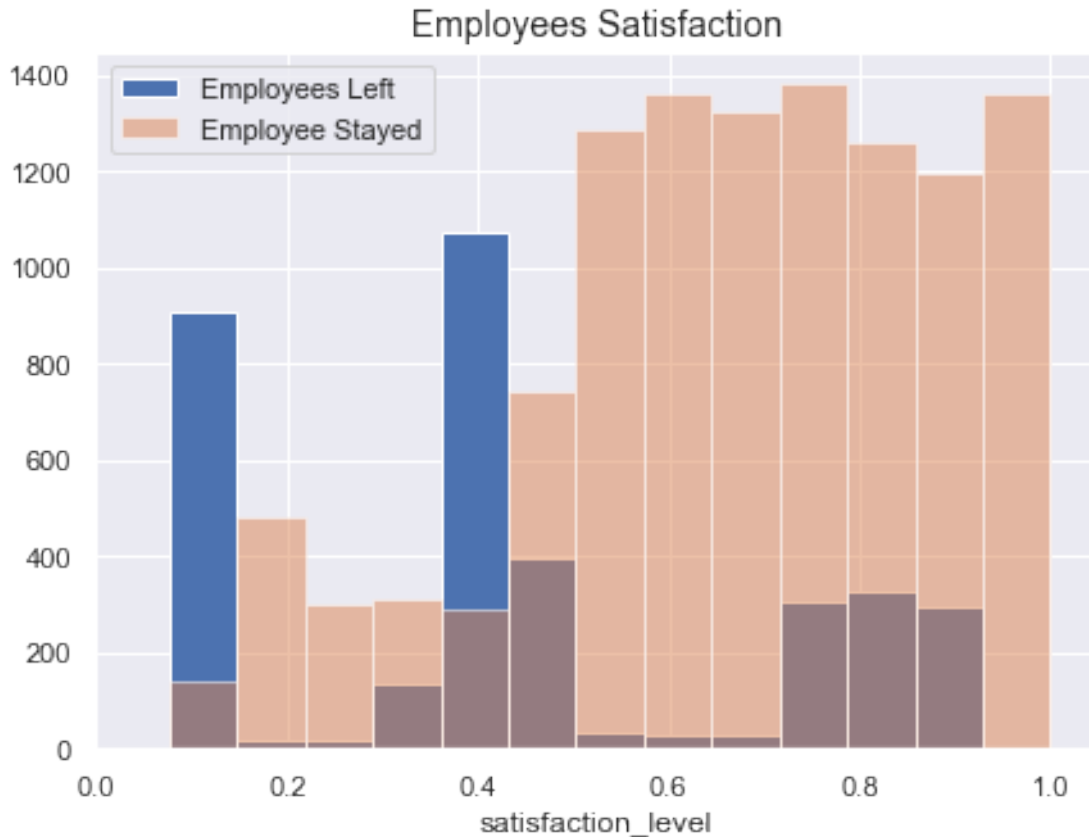
## 0.12 OBSERVATION

- Employees with 7 or more years did not leave.
- In year 1, new hires hardly leave the company, if they do the number is very low. The problem starts when the employees have more than 3 years and get worst when they achieve 5 years.

```
[29]: # Plot comparing employees that left with satisfaction

plt.figure(figsize=(7,5))
bins = np.linspace(0.006,1.000, 15)
plt.hist(hr[hr['left']==1]['satisfaction_level'], bins=bins, alpha=1,
        label='Employees Left')
plt.hist(hr[hr['left']==0]['satisfaction_level'], bins=bins, alpha = 0.5, label=
        'Employee Stayed')
plt.title('Employees Satisfaction', fontsize=14)
plt.xlabel('satisfaction_level')
plt.xlim((0,1.05))
plt.legend(loc='best');
```





[ ]:

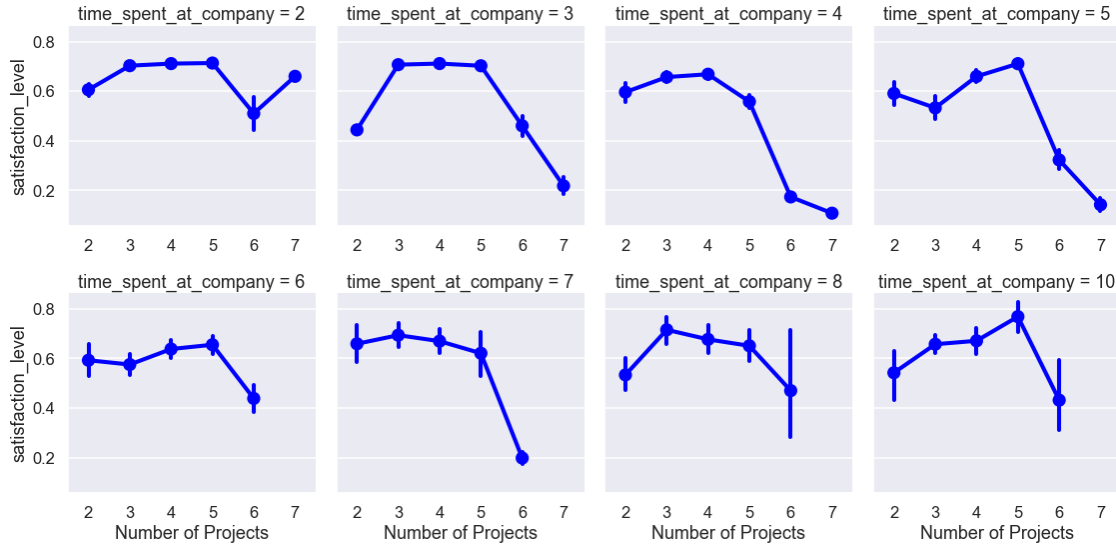
### 0.13 OBSERVATION

- We have a peak at about 0.1 satisfaction level. These are employees who are totally disappointed.
- Another peak at 0.4 satisfaction level, representing another group with the satisfaction level below the average.
- And another amount in the range 0.7 and 0.9, with employees that left, although the high satisfaction.

### 0.14 LINE CHART

```
[45]: sns.set()
sns.set_context("talk")
ax = sns.factorplot(x="number_of_projects", y="satisfaction_level",
    ↳col="time_spent_at_company", col_wrap=4, height=4, color='blue', sharex=False,
    ↳data=hr)
ax.set_xlabels('Number of Projects')
```

[45]: <seaborn.axisgrid.FacetGrid at 0x177315c62c8>



## 0.15 OBSERVATION

It appears clear that there is a drop in satisfaction when employees are working on 6 or more projects.

## 0.16 Summary of the Exploratory Data Analysis

- It is a relatively young company, on average, employees have 3 or 4 years in the company and the oldest employee has been working there for 10 years.
- In five years only 2% of the employees were promoted. Is possible that many employees get unmotivated and start planning to leave.
- Employees with 7 or more years in the company did not left. Employees with 5 years have a greater chance to leaving. Management may have to plan an incentive package to retain employees who work for five years.
- The employees with 4 years in the company have the lowest average satisfaction level of all the company with (0.47).
- The satisfaction drops when the employees are working on 5 or more projects. A number of 3 or 4 projects seems to be ideal independent of the time spent in the company.

## 0.17 MACHINE LEARNING

Now we have to predict who will leave the company

```
[31]: #sklearn libraries needed for machine learning  
  
from sklearn.model_selection import train_test_split
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier

```

```

[32]: #Train-Test split
y = hr['left']
X = hr.drop(['left', 'department'], axis=1)

X.head()

```

```

[32]:      satisfaction_level  last_evaluation  number_of_projects  \
0                0.38                0.53                2
1                0.80                0.86                5
2                0.11                0.88                7
3                0.72                0.87                5
4                0.37                0.52                2

      average_monthly_hours  time_spent_at_company  Work_accident  \
0                157                3                0
1                262                6                0
2                272                4                0
3                223                5                0
4                159                3                0

      promotion_last_5years  salary
0                0            1
1                0            2
2                0            2
3                0            1
4                0            1

```

```

[33]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
      ↪30,random_state=123)

```

```

[34]: from sklearn.preprocessing import StandardScaler
Scaler_X = StandardScaler()
X_train = Scaler_X.fit_transform(X_train)
X_test = Scaler_X.transform(X_test)

```

```

[35]: #Logistic Regression
logis = LogisticRegression()
logis.fit(X_train, y_train)

```

```
logis_score_train = logis.score(X_train, y_train)
print("Training score: ",logis_score_train)
logis_score_test = logis.score(X_test, y_test)
print("Testing score: ",logis_score_test)
```

Training score: 0.7904562339270407  
Testing score: 0.7857777777777778

```
[36]: # Accuracy Matrix
pred=logis.predict(X_test)
print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
```

0.7857777777777778  
[[3183 246]  
 [ 718 353]]

```
[37]: # display table of correlation for each column
coeff_df = pd.DataFrame(X.columns)
coeff_df.columns = ['Features']
coeff_df["Correlation"] = pd.Series(logis.coef_[0])

coeff_df.sort_values(by='Correlation', ascending=False)
```

```
[37]:
```

	Features	Correlation
4	time_spent_at_company	0.370398
3	average_monthly_hours	0.219374
1	last_evaluation	0.157843
6	promotion_last_5years	-0.195991
2	number_of_projects	-0.385212
7	salary	-0.440508
5	Work_accident	-0.554016
0	satisfaction_level	-1.027387

```
[38]: #decision tree
dt = tree.DecisionTreeClassifier()
dt.fit(X_train, y_train)
dt_score_train = dt.score(X_train, y_train)
print("Training score: ",dt_score_train)
dt_score_test = dt.score(X_test, y_test)
print("Testing score: ",dt_score_test)
```

Training score: 1.0  
Testing score: 0.9755555555555555

```
[39]: #decision tree
dt = RandomForestClassifier()
```

```

dt.fit(X_train, y_train)
dt_score_train = dt.score(X_train, y_train)
print("Training score: ",dt_score_train)
dt_score_test = dt.score(X_test, y_test)
print("Testing score: ",dt_score_test)

```

Training score: 0.9978093151728736  
Testing score: 0.9877777777777778

```

[40]: #kNN
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
knn_score_train = knn.score(X_train, y_train)
print("Training score: ",knn_score_train)
knn_score_test = knn.score(X_test, y_test)
print("Testing score: ",knn_score_test)

```

Training score: 0.9723783217449281  
Testing score: 0.9568888888888889

```

[41]: #SVM
svm = SVC()
svm.fit(X_train, y_train)
svm_score_train = svm.score(X_train, y_train)
print("Training score: ",svm_score_train)
svm_score_test = svm.score(X_test, y_test)
print("Testing score: ",svm_score_test)

```

Training score: 0.9665682445947233  
Testing score: 0.9617777777777777

```

[42]: #random forest
rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
rfc_score_train = rfc.score(X_train, y_train)
print("Training score: ",rfc_score_train)
rfc_score_test = rfc.score(X_test, y_test)
print("Testing score: ",rfc_score_test)

```

Training score: 0.9982855510048576  
Testing score: 0.9897777777777778

```

[43]: #Model comparison
models = pd.DataFrame({
    'Model' : ['Logistic Regression', 'SVM', 'kNN', 'Decision_
→Tree', 'Random Forest'],
    'Training_Score' : [logis_score_train, svm_score_train,
→knn_score_train, dt_score_train, rfc_score_train],

```

```

        'Testing_Score' : [logis_score_test, svm_score_test, knn_score_test,
↪dt_score_test, rfc_score_test]
    })
models.sort_values(by='Testing_Score', ascending=False)

```

```

[43]:
      Model  Training_Score  Testing_Score
4  Random Forest      0.998286      0.989778
3  Decision Tree      0.997809      0.987778
1          SVM      0.966568      0.961778
2          kNN      0.972378      0.956889
0 Logistic Regression      0.790456      0.785778

```

## 0.18 Recommendations

- Satisfaction level is the major impact on whether employees stay or leave the company.
- Improve work life balance by having the right number of projects. Employees with 3-4 projects assigned tend to stay. Similarly, number of average hours a month plays a role in employees leaving or staying.
- Provide training so that their evaluation score can improve. The data shows that employees with a low evaluation score are likely to leave.